

KuaiRec: A Fully-observed Dataset for Recommender Systems

Chongming Gao^{1,†}, Shijun Li^{1,†}, Wenqiang Lei^{2,*}, Biao Li³, Peng Jiang³, Jiawei Chen¹,
Xiangnan He¹, Jiaxin Mao⁴, Tat-Seng Chua⁵

¹University of Science and Technology of China; ²Sichuan University, China; ³Kuaishou Technology Co., Ltd.

⁴Renmin University of China; ⁵National University of Singapore

chongming.gao@gmail.com, lishijun@mail.ustc.edu.cn, wenqianglei@gmail.com, biao6@139.com, jp2006@139.com,
cjjwustc@ustc.edu.cn, xiangnanhe@gmail.com, maojiaxin@gmail.com, chuats@comp.nus.edu.sg

ABSTRACT

Recommender systems are usually developed and evaluated on the historical user-item logs. However, most offline recommendation datasets are highly sparse and contain various biases, which hampers the evaluation of recommendation policies. Existing efforts aim to improve the data quality by collecting users' preferences on randomly selected items (e.g., Yahoo! [36] and Coat [47]). However, they still suffer from the high variance issue caused by the sparsely observed data. To fundamentally solve the problem, we present *KuaiRec*, a fully-observed dataset collected from the social video-sharing mobile App, Kuaishou. The feedback of 1,411 users on almost all of the 3,327 videos is explicitly observed. To the best of our knowledge, this is the first real-world fully-observed dataset with millions of user-item interactions in recommendation.

To demonstrate the advantage of *KuaiRec*, we leverage it to explore the key questions in evaluating conversational recommender systems. The experimental results show that two factors in traditional partially-observed data — the data density and the exposure bias — greatly affect the evaluation results. This entails the significance of our fully-observed data in researching many directions in recommender systems, e.g., the unbiased recommendation, interactive/conversational recommendation, and evaluation. We release the dataset and the pipeline implementation for evaluation at <https://chongminggao.github.io/KuaiRec/>.

KEYWORDS

Fully-observed data, Recommendation, Evaluation, User simulation

ACM Reference Format:

Chongming Gao^{1,†}, Shijun Li^{1,†}, Wenqiang Lei^{2,*}, Biao Li³, Peng Jiang³, Jiawei Chen¹, Xiangnan He¹, Jiaxin Mao⁴, Tat-Seng Chua⁵. 2022. *KuaiRec: A Fully-observed Dataset for Recommender Systems*. In . ACM, New York, NY, USA, 11 pages. <https://doi.org/10.1145/1122445.1122456>

[†] Both authors contributed equally to this research. * Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2022 Association for Computing Machinery.
ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00
<https://doi.org/10.1145/1122445.1122456>

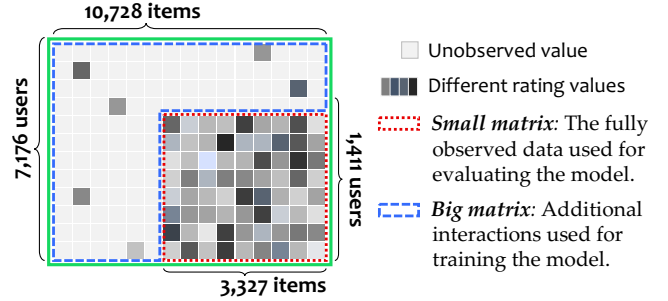


Figure 1: Illustration of the proposed *KuaiRec* dataset.

1 INTRODUCTION

Recommender systems (RSs) are designed to estimate users' preferences on items. A thorny problem in recommendation is how to faithfully evaluate the model, i.e., to tell whether a model can predict the user preference correctly. The most straightforward solution is to conduct the A/B test, which resorts to the online environment to see if real users are satisfied with the recommendations made by the model. However, this is time- and money-consuming and entails the risk of failure [17, 24]. Therefore, it is necessary to evaluate on the offline observed data.

However, the observed data in almost all recommendation datasets is highly sparse, i.e., only a few values are known in the user-item interaction matrix. This imposes great difficulties on evaluation. Specifically, the massive amount of missing values cannot be simply treated as either negative samples or missing-not-at-random (MNAR) data [?]. It means that the sparsely-observed data is not randomly sampled from the whole user-item matrix, which results in the exposure bias permeating through the data generation process in recommendation. The exposure bias can be further divided according to the cause. For example, popularity bias occurs when the model is prone to exposing the popular items [1]; And position bias (or selection bias) is introduced because users have the propensity to select more often the items they like [15, 41]. With these unknown values and pervasive biases in the recommendation data, the results of the offline evaluation are hard to be persuasive, which creates the widely-recognized challenge for evaluating recommender systems [36?].

To remedy the inherent defect of the data, some researchers present recommendation datasets containing randomly sampled data. For example, both the Yahoo! [36] and Coat [47] datasets contain a set of missing-complete-at-random (MCAR) data that can be used for unbiased evaluation. However, these MCAR data of

these datasets is still too sparse, which can bring a high variance to the evaluation results [45].

In this paper, we try to fundamentally address this issue by introducing *KuaiRec*, a fully-observed dataset collected from the social video-sharing mobile App, Kuaishou¹. The term “fully-observed” means there are almost no missing values in the user-item matrix, i.e., each user has viewed each video and left feedback. This fully-observed dataset, dubbed the *small matrix* for convenience, can be used for faithful evaluation. And for the training purpose, we further collect a larger dataset, dubbed the *big matrix*, which contains additional interactions for the users and items in the *small matrix*. Figure 1 gives an illustration of the small matrix and big matrix that make up the *KuaiRec* dataset. Note that all user-item interactions in the *small matrix* are excluded from the *big matrix* to separate the training and evaluation data. Table 1 lists the statistics of *KuaiRec*. Besides the user-item interaction, we further collect the side information of users and items. On the user side, we collected 30 fields representing users’ personal features, which includes 18 encrypted one-hot vectors. We also collect the social relationship among these users. On the item side, we list each item’s tags, i.e., a set of categories such as {Gaming, Sports}. Furthermore, we collect 56 fields to enhance the detailed information of items, which includes 45 fields describing the statistics of each item within each day from July, 5th, 2020 to September 5th, 2020. For example, how many users click, like, or add a video to their favorite list each day.

To the best of our knowledge, *KuaiRec* is the first dataset generated from real-world recommendation logs with all users’ preferences on items known, and its scale (millions of user-item interactions) is much larger than the existing MCAR datasets. This can make the evaluation in offline data as effective as the online A/B test since there is no need to handle the missing values, which can benefit many research directions such as:

- **Unbiased recommendation**, which aims to explicitly identify the biases from data and remove the bad propensity from the recommender [47? ?].
- **Interactive recommendation**, which focuses on improving the recommendation based on the interaction history, i.e., previous recommendations and corresponding user feedback [11? ? ?].
- **Conversational recommender system (CRS)**, which is a kind of interactive recommender system that further utilizes the abundant features of items to efficiently and flexibly identify user preferences [12, 22, 26?].

All of these directions need intensive unbiased data for evaluation. Due to the limited space, we only focus on CRSs in this paper and leave other directions for future research.

To demonstrate the efficacy and advantage of *KuaiRec*, we leverage it to conduct the evaluation of the CRS, which has caught the eyes of the research community recently due to its great potential [12]. Existing solutions for the evaluation of CRSs usually use the user simulation techniques based on the sparsely-observed data. However, the trustworthiness of such approaches remains unknown [12, 64]. With *KuaiRec*, we can explore the research question that cannot be answered in existing studies: *Is partially-observed data as trustworthy as the fully-observed data w.r.t. the evaluation of CRSs?*

Table 1: Statistics of the *KuaiRec* dataset. The density of small matrix is 99.6% instead of 100% because some users have blocked the videos from certain authors. All friends in the social network are in the 7,176 users of the big matrix.

	#users	#Items	#Interactions	Density
<i>Small matrix</i>	1,411	3,327	4,676,570	99.6%
<i>Big matrix</i>	7,176	10,728	12,530,806	16.3%
User feature:	Each user has 30 features which includes 12 explicit features and 18 encrypted vectors.			
Item feature:	Each video has at least 1 and at most 4 tags out of the totally 31 tags, e.g., {Sports}. Each item has 56 explicit features, where 45 fields are the statistics of each day.			
Social network:	<i>Small matrix</i> : 146 users have friends. <i>Big matrix</i> : 472 users have friends.			

As the first attempt to explore the question, we examine two essential factors in partially observed data: data density and exposure bias. The **data density** means the ratio of the items exposed to the users, i.e., the proportion of observed values in the user-item matrix. For the **exposure bias**, we explore three exposure strategies, namely **uniformly random exposure** (which is unbiased), **positivity-oriented exposure** (making exposed user-item interactions contain more historically positive samples to simulate the positivity bias [1]), and **popularity-oriented exposure** (making the exposure lean towards the popular samples to simulate the popularity bias [15, 41]). We can examine the effect of the two factors by sampling part of the user-item interactions from the small matrix as the test set to conduct evaluations. The experimental results on *KuaiRec* provide two key insights: 1) Exposure biases caused by the exposure strategy greatly affect the performances and rankings of different models in evaluation. 2) Even under the uniformly random exposure strategy (i.e., without exposure biases), different data densities can still result in inconsistent results. These insights entail the significance of our fully-observed dataset.

We summarize our contributions as follows:

- We are the first to present a real-world fully-observed dataset (density: almost 100%) in recommendation. It contains millions of dense interactions and rich side information.
- With this dataset, we explore the important questions that cannot be answered in existing works of evaluation.
- The experiments provide interesting insights for evaluating conversational recommender systems. These insights emphasize the importance of the fully-observed dataset.

We release this dataset as well as the scripts for loading and analyzing its statistics at <https://chongminggao.github.io/KuaiRec/>. We also release the pipeline implement for evaluating conversational recommendation at https://github.com/xiwenchao/fully_observed_demo, with the hope to illustrate the advantage of this dataset and support further discussions along with related research topics.

¹<https://www.kuaishou.com/cn>

2 RELATED WORK

In this section, we briefly review the methods, datasets, and problems in offline evaluation of recommender systems. Then, we introduce conversational recommender systems.

2.1 Offline Evaluation in Recommendation

Online A/B tests have become ubiquitous in tech companies for the purpose of assessing the performances of models and rolling out the improved recommender system [14]. However, it usually consumes much time and money and thus is impractical for academic researchers to conduct the evaluation online.

Therefore, researchers usually resort to offline computing indicators in offline data, e.g., Precision, Recall, NDCG [19], and MAP [3]. However, such evaluations suffer from strong assumptions, such as independence between items and user feedback can be translated into a supervised task [14, 37]. This is inconsistent with the nature of the recommender system, which should be treated as a sequential decision problem [45, 56, 58, 73].

To address this, researchers propose two branches of solutions based on the offline data: 1) off-policy evaluation (OPE), which aims to estimate the performance of the target policy using data generated by a different policy [20, 51]; 2) user simulation, whose core idea lies in filling the missing values before the evaluation [16, 64]. However, the former solution suffers from the high variance issue [45], while the latter will inevitably introduce estimation error. The primary cause of the problems is that the offline data is too sparse. In this way, fully-observed data is entailed for exploring how much the sparsity affects the results of the evaluation.

Up to now, far too little attention has been paid to collecting such a dataset. We briefly introduce classic existing high-quality datasets that contain unbiased randomly-sampled data used for offline evaluation in recommendation.

- **Yahoo!** [36]. It contains the conventional missing-not-at-random (MNAR) data that contains approximately 300,000 user-supplied ratings from 15,400 users on 1,000 items in total. It also contains a set of missing-complete-at-random (MCAR) data by asking 5,400 users to give ratings on 10 items that are randomly selected from 1,000 items.
- **Coat** [47]. It collects the ratings of 290 users on 24 self-selected items and 16 randomly-selected items from total 300 items.
- **Open Bandit Dataset** [45]. It contains interactions collected from two logged bandit policies: a Bernoulli Thompson Sampling policy and a uniform (random) policy. There are approximately 26 million interactions collected on users' clicks on 80 items.

However, these datasets are still highly sparse, e.g., Yahoo! has only 54,000 randomly-selected interactions out of $5,400 \times 1,000$ user-item pairs (i.e., density: 1%). So far, there is no dataset containing dense interaction data, not to mention the fully-observed data, to conduct the faithful evaluation in recommendation.

2.2 Conversational Recommender Systems

Recommender systems are powerful tools to help users reach their desirable items from a massive amount of items. However, the traditional static recommendation systems have limitations regarding




Figure 2: Interactions between the CRS and the simulated/real user on the model development/serving stage.

capturing precious real-time user preferences [55, 67] and interpreting user motivation [13, 35, 63]. Therefore, researchers overcome these problems by developing models that work in an online manner, i.e., interactive recommendation systems, such as the multi-armed bandit (MAB) [8, 26, 54, 71] and deep reinforcement learning (DRL) based models [4, 11, 57, 70, 72].

However, most interactive recommendation methods aim to collect user feedback on all recommended items, which suffer from low efficiency as there are too many of them. Critiquing-based methods, as an early form of conversational recommender systems (CRSs), ask users questions about attributes to narrow down the item candidate space [5, 32, 33]. However, there is still a limitation: the models keep alternating asking and recommending, which should be avoided as the recommendation should only be made when the confidence is high. Recently, the emergence of CRSs has solved the problem. Compared with previous methods, a CRS model has an additional user interface that allows for more flexible forms of interaction, e.g., in the form of tags [7, 8], template utterances [10, 21, 22, 26, 68], or free natural language [6, 25, 27, 30, 31, 42, 43, 74–76]. More importantly, a CRS model has a conversation strategy module that controls the core logic of the multi-turn interaction [21–23, 26, 31, 59, 65]. Nonetheless, there are still a lot of challenges in CRSs needed to be addressed [12].

One of the main challenges is the evaluation of CRSs. As discussed in Section 1, researchers have to build user simulators to evaluate CRSs before serving online users (Figure 2). There are typically four types of strategies for simulating users [12]. (1) Directly using the historical user-item interactions as positive samples [4, 21, 22, 50, 77]. If the items recommended by a CRS are in the users' test set, then this recommendation is deemed to be a successful one. (2) Extracting user simulators from user reviews [66, 75]. Unlike the consumption history, the review data of an item usually explicitly mention attributes, which can reflect the personalized opinions of the user on this item. (3) Building user simulators by imitating human conversational corpora [25, 31, 52]. They are mostly used in the dialogue system-driven CRSs, and the idea is to make the deep language model learn the patterns in human conversations or simulated corpora. (4) Estimating user preferences on the unobserved values before the test [8, 65]. All of these user simulators are built by estimating users' preferences and imitating real users' behaviors. However, they are typically learned on the sparsely-observed data that usually contains various biases. [12, 15] point out that the evaluation of CRSs by these biased user simulators is not trustworthy. Thereby, the fully-observed dataset proposed in this work can serve as a necessary testbed for evaluating CRSs.

3 DATA COLLECTION

In this section, we introduce the data collection process and show the representativeness of the collected data.

3.1 The KuaiRec Dataset

We collect the data from Kuaishou App, a famous short-video platform with billions of users. On this platform, users can watch a variety of short-form videos from genres such as dance, entertainment, and fitness/sports. The videos are organized by recommending streaming where each time the user can see only one video. The user can swipe up or down to skip to the last or next videos at any time without having to wait for the end of the video. In this paper, we sometimes refer to videos as *items* for ease of understanding.

3.1.1 Collecting Fully-observed Data (i.e., The Small Matrix). All the users and videos in our dataset, as well as their interaction records, are collected in the period from July, 5th, 2020 to September 5th, 2020 on Kuaishou App. Over this period, we first collect all the users and videos labeled with "high quality" by the data analysis team of Kuai, which are thought to be very valuable for intensive study. Then we further collect a subset of users and videos from them, aiming to enable as many selected videos as possible to have been watched by the selected users. For the missing values, i.e., the rest videos that the user did not watch, we manipulate the online recommendation rule to insert these videos into the recommendation streaming to make sure all of them are exposed to the user. It takes 15 days for this exposure process.

Finally, there are 1,411 users and 3,327 videos that compose a matrix where each element represents a user's feedback on a video. The density of this matrix is 99.6% instead of 100% because some users have explicitly blocked the videos from certain authors. We further conduct the two-sample Kolmogorov-Smirnov tests [38] to demonstrate the *representativeness* of the selected users and videos. To our knowledge, this is the first real-world fully-observed dataset that the ground-truth preference of every user on every item is known. We refer to this fully-observed matrix as the *small matrix*.

3.1.2 The Big Matrix: Peripheral Data of the Small Matrix. The fully-observed dataset can be used for the trustworthy evaluation of CRSs. However, we need additional data for the training of CRSs [21, 22, 50]. Hence, we collect a partially-observed data that contains the interactions of more users and videos. We refer to this larger data as the *big matrix*. It contains 7,176 users and 10,728 videos that include all the users and videos in the *small matrix*. Note that all user-item interactions in the *small matrix* are excluded from the *big matrix* to separate the training and evaluation data.

3.1.3 Defining the Positive Interactions. Our collected raw data contains various user feedback on each watched video such as comments, likes, the total time spent on the video, etc. For simplicity, we follow the conventional practice in the recommender of the Kuaishou App to binarize each user-video pair into either positive or negative samples accordingly to the watch time. Specifically, we define a user-video pair as a positive sample if the user's cumulative watching time of this video is greater than twice the duration of this video, i.e., the user has watched this video completely at least twice. This is deemed as a strong indicator of the users' preference for the videos. The process is illustrated in Figure 3.

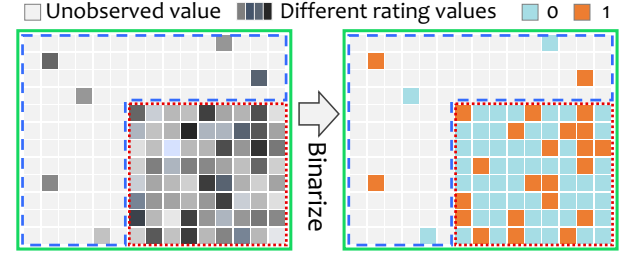


Figure 3: Transformation from the rating matrix to the binary matrix to define the positive samples.

3.1.4 Side Information Collection. We collect the rich features for the users and videos in the *big matrix* as shown in Table 1. Armed with hundreds of these features, researchers can study more comprehensive topics using this dataset.

Especially, we collect 31 genres (i.e., tags) of videos. These tags are manually defined and annotated by the internal annotation team of Kuaishou. Each video is related to at least one and at most four tags. These tags, such as fitness/sports, recipes, are very useful in the question asking and candidate reduction in CRSs models.

Besides, people's minds can be affected by their friends. Therefore, we collect the social network of the 7,176 users in the *big matrix* to benefit the social recommendation [62, 63].

We release the KuaiRec dataset as well as its description and statistics at <https://chongminggao.github.io/KuaiRec/>².

3.2 Hypothesis Testing of Representativeness

Now, we show that the users and videos in the collected the fully-observed *small matrix* are representative of all users and videos in the whole Kuaishou platform. Specifically, we first select six key features (see Table 2) for users and videos, respectively. Then, We conduct the discrete two-sample Kolmogorov-Smirnov (KS) tests [38] to verify whether the users (or videos) in the *small matrix* have the same distribution, w.r.t. the six features, as all users (or videos) in the one-month data of Kuaishou.

Our null hypothesis is that the two distributions have *no significant difference* w.r.t. each feature, i.e., samples from the two data are drawn from the same distribution. We conduct the tests under a 5% rejection level and report the corresponding *p-values* in Table 2.

It's clear to see that the *p-values* of all KS tests on each character are above the rejection level of 5%. It means that we cannot reject the null hypothesis, i.e., users and videos in the fully-observed data are sufficiently representative to reflect the general properties of the users and videos on the whole Kuaishou platform.

4 ENVIRONMENT SETTING FOR CRS

To better illustrate how to utilize the proposed dataset in evaluating CRSs, we focus on the multi-round conversational recommendation (MCR) scenario [21]. We will start with the introduction of MCR as well as how the user simulator is built. Then we introduce how we synthesize the biased partially-observed data to explore its effect on the evaluation.

²The dataset is licensed under CC BY-SA 4.0.

Table 2: The p -values of two-sample KS tests. The selected features for users are {age, location, register channel, operation system, 1st purchase channel, 2nd purchase channel}. For videos, the features are {background music class, visible state, client setting, sound track class, brand class, platform class}. The six features are denoted by Feat1–Feat6.

Feature	Feat1	Feat2	Feat3	Feat4	Feat5	Feat6
User	92.89%	88.27%	6.03%	53.44%	19.00%	6.56%
Video	88.27%	97.62%	99.96%	31.97%	79.74%	99.96%

4.1 Multi-round Conversational Setting

The MCR setting is thought to be the most realistic setting in research so far and has been widely employed in many CRS works [21, 22, 26, 59, 66].

MCR is a typical form-based [18] setting, which means users interact with the system in pre-defined manners instead of dynamically generated natural languages. In MCR, a multi-round interaction process between a CRS model and a user (usually a simulated one) is called one conversation (a.k.a. one conversation session). The system can choose to either ask a question or recommend items to the user at each conversation turn. If the current turn is to recommend, the CRS model recommends top- K items. Otherwise, the CRS model asks a question about attributes such as “Would you like videos with the fitness/sport tag?” Then the system will receive feedback from the user. If the user gives negative feedback to either the question of attributes or recommendation, the system removes them from the candidate pool and continues the conversation. If the user gives positive feedback to the asked attribute, the agent removes all items that do not contain the preferred attribute from the candidate pool. If the user accepts the recommended items, the conversation ends successfully. To best satisfy the users, the objective of CRSs in the MCR setting is to make successful recommendations with the fewest conversation turns.

As discussed in Section 1, researchers built user simulators to automate the evaluation of CRSs. To simulate a user, the simulator needs to give feedback on the items or attributes based on the user preferences estimated from the existing observed data [21, 22, 26, 65, 66]. However, it remains unknown how sparsity and biases in the partially-observed data affect the evaluation.

4.2 Synthesizing Partially-observed Data

Although we can conduct a trustworthy evaluation using the fully-observed data, we want to find out how partially-observed data affect the evaluation. For example, how does the density of the partial data affect the evaluation? And, how do the biases introduced by exposure strategies in the partial-observed data affect the performances and rankings of different models? Therefore, we employ three types of widely used exposure strategies to synthesize the partially-observed data based on our *small matrix* as test sets for evaluation.

4.2.1 Uniformly Random Exposure. Uniformly random is an ideal way to get the data without exposure biases. The data is treated to be a reasonable substitute for fully-observed data. We follow

uniformly random distribution to sample the elements in the *small matrix* without replacement for 9 times, with the density setting to {10%, 20%, ..., 90%} accordingly. The 9 sampled missing at random (MAR) data are used in our experiments to explore how the data density affects the evaluation of CRSs.

4.2.2 Positivity-oriented Exposure. In reality, the partially-observed data is usually biased towards the items that users like, which can be explained for two reasons. First, users tend to choose the items they like to consume or rate, a.k.a., selection bias [37]. Second, the recommendation result can be dominated by the highly-confident items provided by the model [29]. Therefore, we employ positivity-oriented exposure to simulate such biased partially-observed data. Specifically, for each user in the *small matrix*, we sample without replacement a certain number of items with respect to the positivity distribution calculated on the *big matrix*, e.g., if item A is liked by 6 users and item B is liked by 2 users in the *big matrix*, the exposure probability of A should be three times of B. Similarly, with varying the density in {10%, 20%, ..., 90%}, we get 9 sampled missing not at random (MNAR) datasets containing the positivity bias.

4.2.3 Popularity-oriented Exposure. The recommendation model can also bias toward the popular items. Therefore, it is common for the partially-observed data to contain popularity bias [1]. To satisfy the long-tail distribution requirement³, we employ Zipf’s law [39] to decide the probability of each item to be exposed. Specifically, for each user in the *small matrix*, we set the probability of exposing the k -th popular item as $P(k) = \frac{1/k^s}{\sum_{n=1}^N (1/n^s)}$, where $s = 0.5$, $N = 3,327$ (the number of items in the *small matrix*) in our paper. These numbers are set by matching the distributions of popularity in the *big matrix*. This leads the data to follow the long-tail distribution, while the original order of the popularity for items is preserved. Again, we derive the 9 sampled MNAR datasets containing popularity bias with the data density equal to {10%, 20%, ..., 90%}.

5 EXPERIMENTS

In this section, we conduct intensive experiments on the evaluation of CRSs on the KuaiRec dataset. The implementation can be found at https://github.com/xiwenchao/fully_observed_demo.

5.1 Experimental Setting

5.1.1 Research Questions. we will organize our experiments and analyze the results based on the two questions:

RQ1: How does the partially observed data (with and without biases) affect the evaluation of CRSs?

RQ2: Can we improve the evaluation on the partially observed data by estimating the missing values (i.e., matrix completion)?

5.1.2 Evaluation Setting. We conduct experiments on the MCR setting, which has been introduced in Section 4.1 in detail. We follow the configurations of this work [21]. Typically, we set the maximum round of each conversation to 15, since users always have limited patience. And the length of the recommendation list each round is set to 10. To explore the helpfulness of estimating the missing values, we compare the evaluation results of CRSs on different partially observed data before and after matrix completion.

³Real-world recommendation data always satisfies a long-tail distribution [40, 61]

5.1.3 User Simulation. As discussed in Section 1, we need a user simulator to interact with the CRS. In our experiments, we simulate users from the partially-exposed matrix created on the small matrix.

It is worth mentioning that most existing studies [21, 22, 50] used the single-target ground truth (STG) setting to evaluate CRSs. In the STG setting, the target of each conversation is to recommend one ground-truth item. If the recommended top- K items contain the ground-truth item, the user simulator accepts the recommendation and ends the conversation successfully. When the CRS model asks a question, the user simulator gives positive feedback only if the asked attribute belongs to the ground-truth item.

However, in real practice, users tend to have multiple preferences and diverse interests [2, 34]. For example, in the video recommendation streaming, the user may like and accept more than one video recommended by the system. In this case, many videos liked by the user will be mistakenly treated as negative samples under the STG setting. Therefore, we pay more attention to the multi-target ground truth (MTG) setting where users have multiple ground-truth items in a session. Specifically, the user simulator accepts the recommendation only if one of the ground-truth items is in the recommended top- K items. When the CRS model asks a question, the user simulator gives positive feedback if at least one of the ground-truth items contains the queried attribute.

In the evaluation, the number of conversations equals the number of positive user-item pairs in the test set, i.e., we will conduct one conversation for each positive sample in both the STG and MTG settings. The difference is that in the MTG setting we will remove the accepted item from the ground-truth set after a successful recommendation for each simulated user.

5.1.4 Metrics. Following [21, 22, 50], we use two metrics to assess the performance of CRS models in the multi-round conversation. The first is the Average Turns (AT) of all conversations. It is expected to be as small as possible because the model should make successful recommendations with the fewest rounds. Another metric is Success Rate (SR@ t): the proportion of the conversations ended before (or at) round t in all the conversations.

5.1.5 Baselines. We select four representative CRS as baselines. All of them have a recommendation model to measure the similarity between users and items, as well as a conversation strategy to decide whether to ask attributes or make recommendations. We use the factorization machine (FM) model [44] as the recommendation engine in all four methods for ease of comparison. As for the conversation strategy and other configurations, we follow the original paper exactly.

- **Max Entropy** [50]. When asking questions, it always chooses an attribute with the maximum entropy within the current candidate item set.
- **Abs Greedy** [8]. This method focuses on recommending items in every round without asking any questions. It keeps updating the model by treating the rejected items as negative examples.
- **CRM** [50]. It records user preference in a belief tracker and uses reinforcement learning (RL) to find the optimal policy.

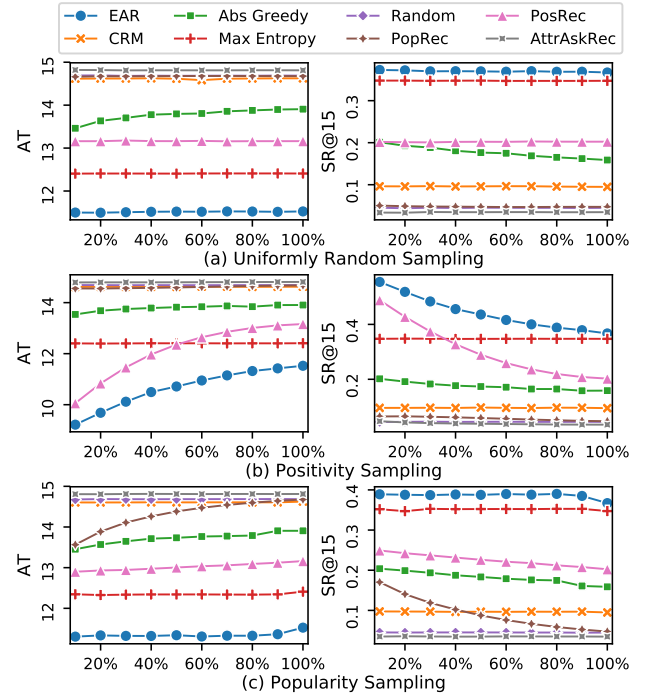


Figure 4: Performance of eight methods on the STG setting by varying the data density and exposure strategies.

- **EAR** [21] A classic CRS model that contains a strategy module based on the RL model similar to CRM, except it considers a more sophisticated state in RL.

Since all the above models contain relatively complex components, e.g., the FM model, we further add four heuristic methods. Each method only contains a naive conversation strategy. This is for better illustrating how biases introduced affect different strategies in evaluation. The four naive methods are:

- **Random.** It randomly selects 10 items from the candidate set to recommend in every turn without asking any questions.
- **Popularity-oriented Recommender (PopRec).** Recommending the top 10 popular items in the candidate sets at every turn without asking any questions.
- **Positivity-oriented Recommender (PosRec).** Similar to PopRec, except it recommends the most positive items in the rest of the item candidates. The most positive items are the items loved by as many users as possible in big matrix.
- **Attribute-asking Recommender (AttrAskRec).** It alternates asking attributes and recommending items. When asking a question, it asks about the most popular attribute; when recommending, it randomly recommends 10 items.

5.2 Exploring the Effects of Partially Observed Data in CRS Evaluation

We use the sampled datasets to explore how the density and exposure strategy in the partially observed data affect the evaluation of CRSs. Specifically, we repeat the evaluation 10 times and report

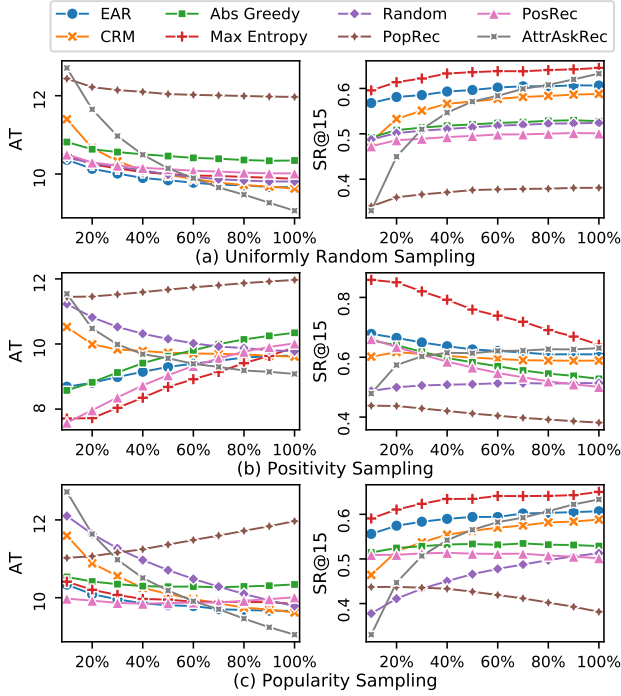


Figure 5: Performance of eight methods on the MTG setting by varying the data density and exposure strategy.

the average AT and average SR@15 over all conversations. The partially-observed data are generated from the *small matrix* following the rules in Section 4.2. It should be noted that all CRS methods are already trained on the *big matrix* before the evaluation.

Instead of paying attention to the absolute performance of the eight methods, we focus on how the rankings (i.e., the relative order) of these CRSs change w.r.t. various densities and exposure strategies. We start with the STG setting and then make the counterpart experiments in the MTG setting with the same configurations.

5.2.1 Analyzing the Results under the STG Setting. We illustrate the results of eight methods on the STG setting under three exposure strategies and 10 densities (including the fully observed data) in Figure 4. From the results, we can get the following insights:

Uniformly Random Exposed Data in Evaluation. In the results of uniformly random exposure, the performance rankings of eight methods keep unchanged and the performances almost keep stable (except for the *Abs Greedy*) over all densities. This is intuitive as the positive samples under the uniformly random exposure are independent and identically distributed. Thus the average performance will not vary much with different data densities due to the law of large numbers [9].

Effects of Positivity Biases in Evaluation. In the results of positivity-oriented exposure, *PosRec* performs satisfyingly at the beginning and becomes worse as the density increases. This is due to its recommendation mechanism (i.e., always recommending the most positive items in the *big matrix*) is correlated with the mechanism of the exposure strategy. Specifically, when the data

density is low, the exposed positive samples are the most positive items, therefore the CRS can find each of them easily in a few turns, hence the AT is low and SR@15 is high. However, when the density increases, more items with a relatively low positivity ratio will be sampled in our synthesized testing dataset. Those items have low rankings in the recommended list of *PosRec*. Therefore the performance deteriorates because of these “hard to predict” items. It is interesting that *EAR* shows a similar trend but with a better performance. It demonstrates *EAR* tends to recommend the positive items but the RL-based conversation strategy makes it perform better. Aside from *EAR* and *PosRec*, other methods have relatively stable absolute performance and consistent rankings, indicating that these methods are not vulnerable to the positivity bias in the partially observed data under the STG setting.

Effects of Popularity Biases in Evaluation. The results of the popularity-oriented exposure have a similar phenomenon: only *PopRec* has unstable performance and inconsistent rankings. It is because *PopRec* always recommends the most popular items, making it easy to find the correct items when the exposed items are the most popular ones. When the unpopular items are exposed with the data density increases, the average performance of *PopRec* is held back by the unpopular items which are hard to predict. Other methods have stable performance and consistent rankings, indicating that they are not sensitive to the popularity bias under the STG setting.

In the STG setting, each item is evaluated independently, which is straightforward for the learning and evaluation of CRSs. However, the users in most real-world applications have multi interests and there is more than one correct answer in the evaluation. Therefore, we explore how things work in the MTG setting.

5.2.2 Analyzing the Results under the MTG Setting. In the MTG setting, we implement the counterpart experiments above, we show the results in Figure 5. It’s obvious that the results are quite different compared to the results in the STG setting. Even under the uniformly random exposure, both the absolute performance and rankings of the CRSs vary a lot with different configurations.

Challenges and Opportunities under the MTG Setting. The instability of the results comes from the complexity of the MTG setting. As discussed in Section 5.1.3, a successful recommendation is achieved when at least one of the ground-truth items is contained in the recommended list. However, multiple ground-truth items do not mean that the recommendation task becomes easier. It is because a system needs to estimate user preferences based on diverse feedback (recall that user simulators give feedback based on multiple preferred items). Under the MTG setting, *Max Entropy*, with a simple conversation strategy, surprisingly performs best on SR@15. The reason that *Max Entropy* outperforms *EAR* might be that *EAR* is specifically designed under the STG setting. However, there is no CRS considering the MTG setting until now. It remains to be a promising direction to develop the CRS under the MTG setting since MTG is more realistic in real life.

Effects of the Question Asking in the MTG Setting. Since the MTG setting is more complicated, the next question is how a model can perform well in this setting. From the results, we find some

clues: we observe that asking attributes affects the performance in the MTG setting more significantly than in the STG setting.

There are two empirical observations: First, the *AttrAskRec* method performs worst in the STG setting, but it performs very well when the data density is large. Second, *CRM* outperforms *Max Entropy* on AT at a certain point when the density increases. We found that *CRM* tends to ask more questions than *Max Entropy*. Specifically, we calculate the average probability for *CRM* to ask a question in each turn is 81.7%, while *Max Entropy* is 66.5%.

To understand the reason, we revisit the mechanism of the user simulator in the MTG setting. When being asked a question about an attribute, the user simulator will return positive feedback if one of the ground-truth items contains the queried attribute. Therefore, when there are few ground-truth items (when the density is small), many questions asked by the CRS will be rejected, which wastes the interaction turn. Thus, the average turns will be high. Conversely, when there are more ground-truth items (when the density is high), it will be easier for the CRS to ask questions to get positive feedback. The positive response can help the CRS quickly narrow down the list of candidate items, which results in small average turns.

Effects of Biases in the MTG Setting. Similar to the observation in the STG setting, the positivity biases and popularity biases can still affect the evaluation results of CRSs. More precisely speaking, the rankings of CRSs change more sharply. For instance, *Random* behaves badly when the density is low on data exposed with positivity bias. In this case, almost all exposed items have a high positivity ratio, so randomly recommending items can be a poor strategy. When the density increases and more items with a relatively low positivity ratio are exposed, the relative performance of *Random* increases. By contrast, the absolute performance of *PosRec* decreases as the density increases. *PosRec* only recommends items with maximal positivity ratio, hence it is unable to recommend items with low positivity ratio. Nevertheless, when the data is sampled uniformly randomly, the variation of absolute performance of both *Random* and *PosRec* is much milder. Therefore, the bias brought by the exposure strategy can aggravate the discordance of rankings.

5.3 Effects of Estimating the Missing Values

After exploring how the density and the exposure strategy in the partially observed data affect the evaluation of CRSs, we further investigate the effects of the remedy: estimating the missing values in the user-item matrix and using the estimated positive samples for building the user simulator in evaluation.

5.3.1 Analyses of Results of Matrix Completion. Estimating missing values, i.e., matrix completion is a well-studied research topic. Due to the limitation of space, we only select three methods related to our research: probabilistic matrix factorization (PMF) [46], exposure matrix factorization (ExpoMF) [28] and its variant ExpoMF-cov [28]. All of these methods are classic baselines of the probabilistic methods in matrix factorization (MF). PMF is an ordinary MF method without considering biases in the data, while ExpoMF and ExpoMF-cov are more sophisticated methods since they use a random variable to model whether a user-item interaction happens, i.e., they can differentiate the unobserved event from a negative preference, thus it shows an advantage in situations where there

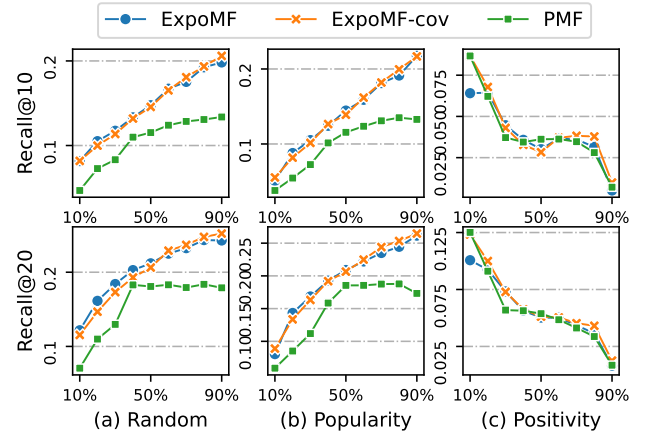


Figure 6: Performance of the matrix completion on the partially-observed data. The training set is the exposed data, i.e., the observed data, and the test set is the unobserved data.

are few exposed interactions. Besides, by modeling the exposure of the events, these two methods can also alleviate biases presented in the partially observed data. The difference between ExpoMF-cov and ExpoMF is that ExpoMF-cov adds the additional exposure covariates to take into account the additional item features.

We train the three methods on the observed part of all datasets with the density and exposure strategy varied and test the performance on the unobserved part, i.e., unsampled values. We report the Recall@10 and Recall@20 on all sampled datasets and illustrate the results in Figure 6. From the results, we observe the following insights:

Effects of Modeling the Exposure Event in Matrix Completion. PMF shows an inferior performance to two debiased methods (ExpoMF and ExpoMF-cov), especially under uniformly random exposure and popularity-oriented exposure. It is due to the fact that ExpoMF and ExpoMF-cov take into account whether the value 0 in the training data represents an unobserved event or a negative sample. Therefore, with the increase of sampled positive values 1, their performances are constantly improved. The curve of PMF shows a clear upward trend when the density is low ($\leq 50\%$), and the growth slows down and stagnates when the density is greater than 50%. It even shows a downward trend in the popularity-oriented exposure when the density is greater than 80%. It means the bottleneck is the model itself and increasing exposed samples cannot further boost the performance. Therefore, distinguishing the unobserved event from negative preference is important.

Effects of Biases in Matrix Completion. The curves of all three methods show a downward trend in the results of the positivity-oriented exposure with the density increasing. It is inevitable because the number of positive samples (i.e., ground-truth preferences) needed to predict decreases sharply with the density increases, and the remaining unexposed positive samples are actually loved by a few users and thus are hard to predict by the model.

Therefore, when the observed data has a larger probability of containing positive samples than the unobserved part, i.e., contains

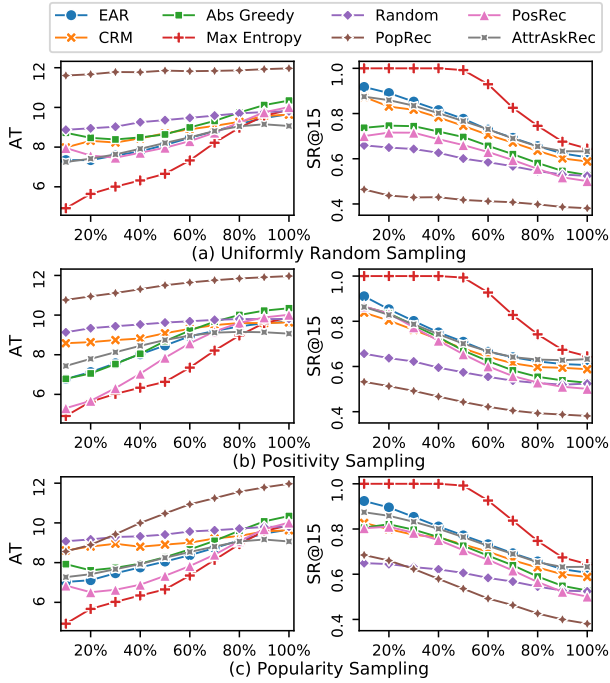


Figure 7: Performance of eight methods on the MTG set using the estimated fully-observed data after completing the missing values via ExpoMF-cov.

positivity bias, the task of completing all missing values in the matrix will be hard. The influence becomes more serious with positivity bias becoming more severe, which explains the downward trend of the curves.

Besides, it’s worth mentioning that the popularity-oriented exposure does not show this phenomenon, because it is independent of popularity that whether an interaction is positive (since the popularity is redefined in our work, see Section 4.2.3). Therefore, the exposed popular data have the same probability of containing positive samples as the unexposed unpopular data, so the MF methods can still learn well from the popular training data and predict satisfyingly on the unpopular test data just like the results under the uniformly random exposure.

5.3.2 Analyses of The Effects of Estimated Values in CRS Evaluation. Due to the limitation of space, we only select ExpoMF-cov, which performs best in our experiments, to estimate the missing values in the partially-observed data and evaluate all CRSs on the completed matrix. Note that the CRSs are pre-trained on the *big matrix*. The results of the CRS evaluation conducted on the completed data are illustrated in Figure 7. We calculate the number of the sampled datasets that have the inconsistent ranking of a CRS before and after matrix completion in Table 3. An inconsistent ranking means the ranking of a CRS on the current test set is inconsistent with the counterpart on the fully observed data. For example, the ranking of *EAR* is 5 on the fully observed data and is 7 on a sampled dataset under the positivity-oriented exposure when the density is 10%, then *EAR* has an inconsistent ranking under the two different datasets.

Table 3: The number of the sampled datasets that have the inconsistent ranking of a CRS *before/after* matrix completion.

Methods	Random		Popularity		Positivity	
	AT	SR@15	AT	SR@15	AT	SR@15
EAR	9/9	7/5	9/9	7/6	9/9	6/5
CRM	9/9	5/0	9/9	4/4	9/9	3/3
Abs Greedy	4/7	3/0	7/8	3/4	6/7	3/3
Max Entropy	5/9	0/0	9/9	0/0	9/9	0/0
Random	5/9	3/8	9/9	9/8	9/9	8/8
PopRec	1/0	1/0	3/2	3/2	1/0	0/2
PosRec	5/9	2/8	9/9	8/8	8/8	8/8
AttrAskRec	6/8	7/5	7/8	7/6	7/8	6/5

Effects of Completing the Partially Observed Data. By analyzing the results w.r.t. SR@15 in Table 3, we find that matrix completion can alleviate the inconsistency of CRSs in some cases. For instance, under the random exposure, *EAR* has inconsistent rankings on 7 sampled partially observed datasets with different data densities. The number reduces to 5 after estimating the missing values. There are nearly half of the values in the table decrease after matrix completion, and the improvement is evident under the random exposure. This validates the helpfulness of matrix completion.

However, the matrix completion does not help restore the true rankings of CRSs w.r.t. AT, where only the number of *PopRec* decreases. We argue that it is because most CRSs have the indistinct absolute performance w.r.t. AT (see Figure 5), and the estimating error introduced by matrix completion can further blur the difference between them. By contrast, the absolute performances of different CRSs on SR@15 are more distinguishable, hence it is easier to restore the truth rankings after completing the matrix, even though the matrix completion introduces some errors.

Effects of Biases after Estimating Missing Values. Generally, even after estimating the missing data, the number of inconsistent rankings of the uniformly randomly sampled data is smaller than the sampled data with biases. This indicates that the biases in the evaluation data persist to affect the CRS evaluation even after estimating the missing values. In some cases, completing the missing values can bring additional errors that make the evaluation worse. Therefore, estimating the missing values in partially observed data can only help identify the true evaluation rankings of CRSs in certain cases where the biases are not serious.

Therefore, these insights further highlight the importance of our fully-observed dataset.

6 CONCLUSION

In this paper, we present KuaiRec, a fully-observed dataset in recommender systems. We focus on using this dataset to study how the data density and exposure bias in the traditional partially-observed data affect the evaluation of CRSs. Extensive experiments on KuaiRec provide interesting insights into the evaluation of CRSs.

However, there are loose ends to our discussions. We hope our unique fully-observed data can support more research in a broader

context. First, it can serve as a testbed for building trustworthy user simulators using partially observed user-item interactions. Although the matrix completion in our experiments demonstrates limited help, it still remains an open question on whether it is possible to use partially observed data to simulate fully-observed data correctly. Our fully-observed data can further support this exploration. Second, the fully-observed dataset can serve as a benchmark dataset for many research directions in recommendation, such as debiasing in recommender systems, interactive recommendation, and evaluation. At least, by releasing this fully-observed data, we want to encourage the efforts of collecting more fully-observed datasets with richer properties, such as multiple domains or more diverse demographics.

REFERENCES

- [1] Himan Abdollahpouri and Masoud Mansoury. 2020. Multi-sided exposure bias in recommendation. *International Workshop on Industrial Recommendation Systems (IRS2020) in Conjunction with ACM KDD '2020* (2020).
- [2] Haris Aziz. 2020. Strategyproof multi-item exchange under single-minded dichotomous preferences. *Autonomous Agents and Multi-Agent Systems* 34, 1 (2020), 1–11.
- [3] Ricardo Baeza-Yates, Berthier Ribeiro-Neto, et al. 1999. *Modern Information Retrieval*. Vol. 463. ACM press New York.
- [4] Haokun Chen, Xinyi Dai, Han Cai, Weinan Zhang, Xuejian Wang, Ruiming Tang, Yuzhou Zhang, and Yong Yu. 2019. Large-Scale Interactive Recommendation with Tree-Structured Policy Gradient. In *AAAI '19*, Vol. 33. 3312–3320.
- [5] Li Chen and Pearl Pu. 2012. Critiquing-based Recommenders: Survey and Emerging Trends. *User Modeling and User-Adapted Interaction* 22, 1–2 (2012), 125–150.
- [6] Qibin Chen, Junyang Lin, Yichang Zhang, Ming Ding, Yukuo Cen, Hongxia Yang, and Jie Tang. 2019. Towards Knowledge-Based Recommender Dialog System. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP '2019)*. 1803–1813.
- [7] Konstantina Christakopoulou, Alex Beutel, Rui Li, Sagar Jain, and Ed H. Chi. 2018. Q&R: A Two-Stage Approach toward Interactive Recommendation. In *KDD '18*. 139–148.
- [8] Konstantina Christakopoulou, Filip Radlinski, and Katja Hofmann. 2016. Towards Conversational Recommender Systems. In *KDD '16*. 815–824.
- [9] Frederik Michel Dekking, Cornelis Kraaikamp, Hendrik Paul Lopuhaä, and Ludolf Erwin Meester. 2005. *A Modern Introduction to Probability and Statistics: Understanding Why and How*. Springer Science & Business Media.
- [10] Yang Deng, Yaliang Li, Fei Sun, Bolin Ding, and Wai Lam. 2021. Unified Conversational Recommendation Policy Learning via Graph-Based Reinforcement Learning (*SIGIR '21*). 1431–1441.
- [11] Chongming Gao, Wenqiang Lei, Jiawei Chen, Shiqi Wang, Xiangnan He, Shijun Li, Biao Li, Yuan Zhang, and Peng Jiang. 2022. CIRS: Bursting Filter Bubbles by Counterfactual Interactive Recommender System. *arXiv preprint arXiv:2204.01266* (2022).
- [12] Chongming Gao, Wenqiang Lei, Xiangnan He, Maarten de Rijke, and Tat-Seng Chua. 2021. Advances and Challenges in Conversational Recommender Systems: A Survey. *AI Open* 2 (2021), 100–126.
- [13] Chongming Gao, Shuai Yuan, Zhong Zhang, Hongzhi Yin, and Junming Shao. 2019. BLOMA: Explain Collaborative Filtering via Boosted Local Rank-One Matrix Approximation. In *DASFAA '19*. Springer, 487–490.
- [14] Alexandre Gilotte, Clément Calauzènes, Thomas Nedelec, Alexandre Abraham, and Simon Dollé. 2018. Offline A/B Testing for Recommender Systems. In *WSDM '18*. 198–206.
- [15] Jin Huang, Harrie Oosterhuis, Maarten de Rijke, and Herke van Hoof. 2020. Keeping Dataset Biases out of the Simulation: A Biased Simulator for Reinforcement Learning Based Recommender Systems. In *RecSys '20*. 190–199.
- [16] Eugene Ie, Chih-wei Hsu, Martin Mladenov, Vihan Jain, Sanmit Narvekar, Jing Wang, Rui Wu, and Craig Boutilier. 2019. Recsim: A Configurable Simulation Platform for Recommender Systems. *arXiv preprint arXiv:1909.04847* (2019).
- [17] Rolf Jagerman, Ilya Markov, and Maarten de Rijke. 2019. When People Change Their Mind: Off-Policy Evaluation in Non-Stationary Recommendation Environments. In *WSDM '19*. 447–455.
- [18] Dietmar Jannach, Ahtsham Manzoor, Wanling Cai, and Li Chen. 2021. A Survey on Conversational Recommender Systems. *ACM Computing Survey* 54, 5, Article 105 (may 2021), 36 pages.
- [19] Kalervo Järvelin and Jaana Kekäläinen. 2000. IR Evaluation Methods for Retrieving Highly Relevant Documents. In *SIGIR '00*. 41–48.
- [20] Damien Lefortier, Adith Swaminathan, Xiaotao Gu, Thorsten Joachims, and Maarten de Rijke. 2016. Large-scale Validation of Counterfactual Learning Methods: A Test-bed. *arXiv preprint arXiv:1612.00367* (2016).
- [21] Wenqiang Lei, Xiangnan He, Yisong Miao, Qingyun Wu, Richang Hong, Min-Yen Kan, and Tat-Seng Chua. 2020. Estimation-Action-Reflection: Towards Deep Interaction Between Conversational and Recommender Systems. In *WSDM '20*. ACM, 304–312.
- [22] Wenqiang Lei, Gangyi Zhang, Xiangnan He, Yisong Miao, Xiang Wang, Liang Chen, and Tat-Seng Chua. 2020. Interactive Path Reasoning on Graph for Conversational Recommendation. In *KDD '20*. 2073–2083.
- [23] Mike Lewis, Denis Yarats, Yann Dauphin, Devi Parikh, and Dhruv Batra. 2017. Deal or No Deal? End-to-End Learning of Negotiation Dialogues. In *EMNLP '17*. 2443–2453.
- [24] Lihong Li, Jin Young Kim, and Imed Zitouni. 2015. Toward Predicting the Outcome of an A/B Experiment for Search Relevance. In *WSDM '15*. 37–46.
- [25] Raymond Li, Samira Ebrahimi Kahou, Hannes Schulz, Vincent Michalski, Laurent Charlin, and Chris Pal. 2018. Towards Deep Conversational Recommendations. In *NeurIPS '18*. 9748–9758.
- [26] Shijun Li, Wenqiang Lei, Qingyun Wu, Xiangnan He, Peng Jiang, and Tat-Seng Chua. 2021. Seamlessly Unifying Attributes and Items: Conversational Recommendation for Cold-Start Users. *TOIS '21* (2021).
- [27] Shuokai Li, Ruobing Xie, Yongchun Zhu, Xiang Ao, Fuzhen Zhuang, and Qing He. 2022. User-Centric Conversational Recommendation with Multi-Aspect User Modeling. *SIGIR* (2022).
- [28] Dawen Liang, Laurent Charlin, James McInerney, and David M. Blei. 2016. Modeling User Exposure in Recommendation. In *WWW '16*. 951–961.
- [29] Dugang Liu, Pengxiang Cheng, Zhenhua Dong, Xiuqiang He, Wei Ke Pan, and Zhong Ming. 2020. A General Knowledge Distillation Framework for Counterfactual Recommendation via Uniform Data. In *SIGIR '20*. 831–840.
- [30] Zeming Liu, Haifeng Wang, Zheng-Yu Niu, Hua Wu, Wanxiang Che. 2021. Durecdial 2.0: A bilingual parallel corpus for conversational recommendation. *EMNLP* (2021).
- [31] Zeming Liu, Haifeng Wang, Zheng-Yu Niu, Hua Wu, Wanxiang Che, and Ting Liu. 2020. Towards Conversational Recommendation over Multi-Type Dialogs. In *ACL '20*. 1036–1049.
- [32] Kai Luo, Scott Sanner, Ga Wu, Hanze Li, and Hoi Jin Yang. 2020. Latent Linear Critiquing for Conversational Recommender Systems. In *WWW '20*. 2535–2541.
- [33] Kai Luo, Hoi Jin Yang, Ga Wu, and Scott Sanner. 2020. Deep Critiquing for VAE-Based Recommender Systems. In *SIGIR '20*. 1269–1278.
- [34] Miltiadis D Lytras, Anna Visvizi, Prasanta Kr Chopdar, Akila Sarirete, and Wade Alhalabi. 2021. Information Management in Smart Cities: Turning end users' views into multi-item scale development, validation, and policy-making recommendations. *International Journal of Information Management* 56 (2021), 102146.
- [35] Jianxin Ma, Chang Zhou, Peng Cui, Hongxia Yang, and Wenwu Zhu. 2019. Learning Disentangled Representations for Recommendation. In *NeurIPS '20*. 5711–5722.
- [36] Benjamin M. Marlin and Richard S. Zemel. 2009. Collaborative Prediction and Ranking with Non-Random Missing Data. In *RecSys '09*. 5–12.
- [37] Benjamin M. Marlin, Richard S. Zemel, Sam Roweis, and Malcolm Slaney. 2007. Collaborative Filtering and the Missing at Random Assumption. In *UAI '07*. 267–275.
- [38] George Marsaglia, Wai Wan Tsang, Jingbo Wang, et al. 2003. Evaluating Kolmogorov's distribution. *Journal of statistical software* 8, 18 (2003), 1–4.
- [39] Mark EJ Newman. 2005. Power laws, Pareto Distributions and Zipf's Law. *Contemporary Physics* 46, 5 (2005), 323–351.
- [40] Yoon-Joo Park and Alexander Tuzhilin. 2008. The Long Tail of Recommender Systems and How to Leverage It. In *RecSys '08*. 11–18.
- [41] Bruno Pradel, Nicolas Usunier, and Patrick Gallinari. 2012. Ranking with Non-Random Missing Ratings: Influence of Popularity and Positivity on Evaluation Metrics. In *RecSys '12*. 147–154.
- [42] Xuhui Ren, Hongzhi Yin, Tong Chen, Hao Wang, Quoc Viet Hung Nguyen, Zi Huang, and Xiangliang Zhang. 2020. CRSAL: Conversational Recommender Systems with Adversarial Learning. *ACM Transactions on Information Systems* 0, ja (2020).
- [43] Zhaochun Ren, Zhi Tian, Dongdong Li, Pengjie Ren, Liu Yang, Xin Xin, Huasheng Liang, Maarten de Rijke, and Zhumin Chen. 2022. Variational Reasoning about User Preferences for Conversational Recommendation. *SIGIR* (2022).
- [44] Steffen Rendle. 2010. Factorization Machines. In *ICDM '10*. 995–1000.
- [45] Yuta Saito, Shunsuke Aihara, Megumi Matsutani, and Yusuke Narita. 2021. Open Bandit Dataset and Pipeline: Towards Realistic and Reproducible Off-Policy Evaluation. In *NeurIPS '21*.
- [46] Ruslan Salakhutdinov and Andriy Mnih. 2007. Probabilistic Matrix Factorization. In *NeurIPS '07*. 1257–1264.
- [47] Tobias Schnabel, Adith Swaminathan, Ashudeep Singh, Navin Chandak, and Thorsten Joachims. 2016. Recommendations as Treatments: Debiasing Learning and Evaluation. In *ICML '16*. 1670–1679.
- [48] Bichen Shi, Makbule Gulcin Osoy, Neil Hurley, Barry Smyth, Elias Z. Trajos, James Geraci, and Aonghus Lawlor. 2019. PyRecGym: A Reinforcement

- Learning Gym for Recommender Systems. In *RecSys '19* (Copenhagen, Denmark). Association for Computing Machinery, New York, NY, USA, 491–495. <https://doi.org/10.1145/3298689.3346981>
- [49] Jing-Cheng Shi, Yang Yu, Qing Da, Shi-Yong Chen, and An-Xiang Zeng. 2019. Virtual-Taobao: Virtualizing Real-World Online Retail Environment for Reinforcement Learning. In *AAAI '19*, Vol. 33. 4902–4909.
- [50] Yueming Sun and Yi Zhang. 2018. Conversational Recommender System. In *SIGIR '18*. 235–244.
- [51] Adith Swaminathan and Thorsten Joachims. 2015. Counterfactual Risk Minimization: Learning from Logged Bandit Feedback. In *ICML '15*. 814–823.
- [52] Svitlana Vakulenko, Evangelos Kanoulas, and Maarten de Rijke. 2020. An Analysis of Mixed Initiative and Collaboration in Information-Seeking Dialogues. In *SIGIR '20*. 2085–2088.
- [53] Ivan Vendrov, Tyler Lu, Qingqing Huang, and Craig Boutilier. 2020. Gradient-based Optimization for Bayesian Preference Elicitation. In *AAAI '20*.
- [54] Qing Wang, Chunqiu Zeng, Wubai Zhou, Tao Li, S Sitharama Iyengar, Larisa Shwartz, and Genady Ya Grabarnik. 2018. Online Interactive Collaborative Filtering Using Multi-Armed Bandit with Dependent Arms. *IEEE Transactions on Knowledge and Data Engineering (TKDE)* 31, 8 (2018), 1569–1580.
- [55] Wenjie Wang, Fuli Feng, Xiangnan He, Hanwang Zhang, and Tat-Seng Chua. 2020. “Click” Is Not Equal to “Like”: Counterfactual Recommendation for Mitigating Clickbait Issue. *arXiv preprint arXiv:2009.09945* (2020).
- [56] Zhenlei Wang, Jingsen Zhang, Hongteng Xu, Xu Chen, Yongfeng Zhang, Wayne Xin Zhao, and Ji-Rong Wen. 2021. Counterfactual Data-Augmented Sequential Recommendation. In *SIGIR '21*. 347–356.
- [57] Yikun Xian, Zuohui Fu, S. Muthukrishnan, Gerard de Melo, and Yongfeng Zhang. 2019. Reinforcement Knowledge Graph Reasoning for Explainable Recommendation. In *SIGIR '19*. 285–294.
- [58] Xin Xin, Alexandros Karatzoglou, Ioannis Arapakis, and Joemon M. Jose. 2020. Self-Supervised Reinforcement Learning for Recommender Systems. In *SIGIR '20*. 931–940.
- [59] Kerui Xu, Jingxuan Yang, Jun Xu, Sheng Gao, Jun Guo, and Ji-Rong Wen. 2021. Adapting User Preference to Online Feedback in Multi-round Conversational Recommendation. In *WSDM '21*.
- [60] Longqi Yang, Yin Cui, Yuan Xuan, Chenyang Wang, Serge Belongie, and Deborah Estrin. 2018. Unbiased Offline Recommender Evaluation for Missing-Not-at-Random Implicit Feedback. In *RecSys '18*. 279–287.
- [61] Hongzhi Yin, Bin Cui, Jing Li, Junjie Yao, and Chen Chen. 2012. Challenging the Long Tail Recommendation. *Proceedings of the VLDB Endowment* 5, 9 (2012), 896–907.
- [62] Junliang Yu, Min Gao, Jundong Li, Hongzhi Yin, and Huan Liu. 2018. Adaptive implicit friends identification over heterogeneous network for social recommendation. In *Proceedings of the 27th ACM International Conference on Information and Knowledge Management*. 357–366.
- [63] Junliang Yu, Min Gao, Hongzhi Yin, Jundong Li, Chongming Gao, and Qinyong Wang. 2019. Generating Reliable Friends via Adversarial Training to Improve Social Recommendation. In *ICDM '19*. IEEE, 768–777.
- [64] Shuo Zhang and Krisztian Balog. 2020. Evaluating Conversational Recommender Systems via User Simulation. In *KDD '20*. 1512–1520.
- [65] Xiaoying Zhang, Hong Xie, Hang Li, and John C.S. Lui. 2020. Conversational Contextual Bandit: Algorithm and Application. In *WWW '20*. 662–672.
- [66] Yongfeng Zhang, Xu Chen, Qingyao Ai, Liu Yang, and W. Bruce Croft. 2018. Towards Conversational Search and Recommendation: System Ask, User Respond. In *CIKM '18*. 177–186.
- [67] Yang Zhang, Fuli Feng, Xiangnan He, Tianxin Wei, Chonggang Song, Guohui Ling, and Yongdong Zhang. 2021. Causal Intervention for Leveraging Popularity Bias in Recommendation. In *SIGIR '21*. 11–20.
- [68] Yiming Zhang, Lingfei Wu, Qi Shen, Yitong Pang, Zhihua Wei, Fangli Xu, Bo Long, and Jian Pei. 2022. Multiple Choice Questions Based Multi-Interest Policy Learning for Conversational Recommendation (*WWW '22*). 2153–2162.
- [69] Xiangyu Zhao, Long Xia, Zhuoye Ding, Dawei Yin, and Jiliang Tang. 2019. Toward Simulating Environments in Reinforcement Learning Based Recommendations. *arXiv preprint arXiv:1906.11462* (2019).
- [70] Xiangyu Zhao, Liang Zhang, Zhuoye Ding, Long Xia, Jiliang Tang, and Dawei Yin. 2018. Recommendations with Negative Feedback via Pairwise Deep Reinforcement Learning. In *KDD '18*. 1040–1048.
- [71] Xiaoxue Zhao, Weinan Zhang, and Jun Wang. 2013. Interactive Collaborative Filtering. In *CIKM '13*. 1411–1420.
- [72] Guanjie Zheng, Fuzheng Zhang, Zihan Zheng, Yang Xiang, Nicholas Jing Yuan, Xing Xie, and Zhenhui Li. 2018. DRN: A Deep Reinforcement Learning Framework for News Recommendation. In *WWW '18*. 167–176.
- [73] Kun Zhou, Hui Wang, Wayne Xin Zhao, Yutao Zhu, Sirui Wang, Fuzheng Zhang, Zhongyuan Wang, and Ji-Rong Wen. 2020. S3-Rec: Self-Supervised Learning for Sequential Recommendation with Mutual Information Maximization. In *CIKM '20*. 1893–1902.
- [74] Kun Zhou, Wayne Xin Zhao, Shuqing Bian, Yuanhang Zhou, Ji-Rong Wen, and Jingsong Yu. 2020. Improving Conversational Recommender Systems via Knowledge Graph based Semantic Fusion. In *SIGKDD '20*. 1006–1014.
- [75] Kun Zhou, Yuanhang Zhou, Wayne Xin Zhao, Xiaoke Wang, and Ji-Rong Wen. 2020. Towards Topic-Guided Conversational Recommender System. In *COLING '2020*.
- [76] Yuanhang Zhou, Kun Zhou, Wayne Xin Zhao, Cheng Wang, Peng Jiang, and He Hu. 2022. C2-CRS: Coarse-to-Fine Contrastive Learning for Conversational Recommender System. *WSDM* (2022).
- [77] Jie Zou, Yifan Chen, and Evangelos Kanoulas. 2020. Towards Question-Based Recommender Systems. In *SIGIR '20*. 881–890.